

Mathematical aspects in machine learning

Lorenz Richter
BTU Cottbus-Senftenberg

November 13, 2017

Contents

1	Statistical learning theory	1
1.1	Concentration inequalities	3
1.2	Error bounds for finite classes	6
1.3	Error bounds for infinite classes	8
2	Bibliography	15

1 Statistical learning theory

We consider data from an input space \mathcal{X} and an output space \mathcal{Y} , specifically the sample $S_n = ((X_i, Y_i))_{i=1}^n \subset \mathcal{X} \times \mathcal{Y}$. The goal is to learn a (prediction) function $f : \mathcal{X} \rightarrow \mathcal{Y}$ that maps input to output data.

Example 1.1. $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \{-1, 1\}$ (*binary classification*), $\mathcal{Y} = \mathbb{N}$, $\mathcal{Y} = \mathbb{R}$ (*regression*).

In order to develop a proper theory we need to make some assumptions:

- There exists an unknown probability distribution \mathbb{P} over $\mathcal{X} \times \mathcal{Y}$.
- The data S_n are i.i.d. from \mathbb{P} , i.e. $(X_i, Y_i) \sim \mathbb{P}$ for every $i = 1, \dots, n$.
- The future data (sometimes called test data) also come from \mathbb{P} .

In order to measure how good we learn the prediction function f we consider a loss function $l : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}^+$ that indicates deviations from predictions and true values. Being stochastically more meaningful, we consider for any function f the expected loss

$$L(f) := \mathbb{E}_{(X,Y) \sim \mathbb{P}}[l(f(X), Y)],$$

which in statistics is sometimes called “(Bayes-)Risk”. The goal is to come up with a learning algorithm $\mathcal{A} : S_n \mapsto \hat{f}_n$ (i.e. $\hat{f}_n := \mathcal{A}(S_n)$, indicating that the function depends on the n training data) s.t. $L(\hat{f}_n)$ is small. We therefore consider $\mathbb{E}[l(\hat{f}_n(X), Y)|S_n]$, i.e. the expectation conditioned on the samples, which is still a random quantity (as it depends on the sample data)¹.

Definition 1.2. A predictor f^B is called *Bayes-optimal* if it minimizes the expected loss, i.e.

$$L(f^B) = \inf_f L(f) =: L^B.$$

Remark 1.3. Bayes-optimality depends on $\mathbb{P}, \mathcal{X} \times \mathcal{Y}$ and l .

Example 1.4 (Zero noise or function learning). One could consider the case that our targets Y are deterministically prescribed by a function g , i.e. $\mathbb{P}(Y = g(X)|X = x) = 1$. However, this is a rather unrealistic case.

¹We will continuously omit the measure \mathbb{P} in the expected value.

Example 1.5 (Binary classification). Consider $\mathcal{Y} = \{0, 1\}$ and the loss $l(y', y) = \mathbb{1}\{y' \neq y\}$. One can show that the Bayes-optimal predictor (in this case classifier) is $f^B(x) = \mathbb{1}\{\eta(x) > \frac{1}{2}\}$ with $\eta(x) = \mathbb{P}(Y = 1|X = x)$. We have

$$\begin{aligned}\mathbb{P}(f(X) \neq Y|X = x) &= 1 - \mathbb{P}(f(X) = Y|X = x) \\ &= 1 - (\mathbb{1}_{\{f(x)=1\}}\eta(x) - \mathbb{1}_{\{f(x)=-1\}}(1 - \eta(x))) .\end{aligned}$$

This yields

$$\begin{aligned}\mathbb{P}(f(X) \neq Y|X = x) - \mathbb{P}(f^B(X) \neq Y|X = x) &= \eta(x) (\mathbb{1}_{\{f^B(x)=1\}} - \mathbb{1}_{\{f(x)=1\}}) + (1 - \eta(x)) (\mathbb{1}_{\{f^B(x)=-1\}} - \mathbb{1}_{\{f(x)=-1\}}) \\ &= (2\eta(x) - 1) (\mathbb{1}_{\{f^B(x)=1\}} - \mathbb{1}_{\{f(x)=1\}}) \geq 0\end{aligned}$$

and by integration w.r.t. x the statement

Example 1.6 (Regression). We consider $\mathcal{Y} = \mathbb{R}$ and $l(y', y) = (y' - y)^2$. Then the Bayes-optimal predictor is $f^B(x) = \mathbb{E}[Y|X = x]$, which we can see by noting that

$$\begin{aligned}L(f) &= \mathbb{E}[(f(X) - Y)^2] \\ &= \mathbb{E}[(f(X) - f^B(X) + f^B(X) - Y)^2] \\ &= \mathbb{E}[(f(X) - f^B(X))^2] + \mathbb{E}[(f^B(X) - Y)^2] - 2\mathbb{E}[(f(X) - f^B(X))(f^B(X) - Y)]\end{aligned}$$

is minimized by $f = f^B$ since the last term is equal to

$$\mathbb{E}_X [\mathbb{E}_{Y|X} [(f(X) - f^B(X))(f^B(X) - Y)]] = \mathbb{E}_X [(f(X) - f^B(X))\mathbb{E}_{Y|X} [(f^B(X) - Y)]] = 0.$$

Let us come back to finding a good function f . Somehow we need to work with our sample S_n . We therefore define the empirical risk to be

$$L_n(f) = \frac{1}{n} \sum_{i=1}^n l(f(X_i), Y_i).$$

We see that $\mathbb{E}[L_n(f)] = L(f)$ and can therefore consider empirical risk minimization (ERM), namely

$$\hat{f}_n := \arg \inf_f L_n(f)$$

as a reasonable strategy to learn the function f . However, the following example shows that this is not always a good idea.

Example 1.7. Consider $\mathcal{X} = [0, 1]$, $\mathcal{Y} = \{0, 1\}$, $l(y', y) = \mathbb{1}\{y' \neq y\}$ and $\mathbb{P}(Y = 1|x) = \frac{1}{2}$ for all $x \in \mathcal{X}$. Define the predictor as

$$\hat{f}_n(x) = \begin{cases} y & \text{if } (x, y) \in S_n \text{ for some value } y \\ 0 & \text{otherwise} \end{cases}.$$

Note again that this predictor is not fixed, but depends on the data. We can see that it performs very badly: $L_n(\hat{f}_n) = 0$, but $L(\hat{f}_n) = \frac{1}{2}$. This is a stereotypical example of what is described as overfitting.

Theorem 1.8 (No-free-lunch). Let \mathcal{A} be any learning algorithm for binary classification and let $n < \frac{|\mathcal{X}|}{2}$ be a sample size. Then there exists a probability distribution \mathbb{P} s.t.

- (i) there exists a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ with $L(f) = 0$,
- (ii) with probability $\geq \frac{1}{7}$ over the sample $S_n \sim \mathbb{P}^n$ we have $L(\mathcal{A}(S_n)) > \frac{1}{8}$.

Proof. See chapter 5 in [7]². □

Remark 1.9. This means that there is no perfect algorithm that works for all distributions. Bayes optimality cannot be obtained independent of \mathbb{P} .

²The proof is similar to the concept of VC dimension, which we will discuss later.

In order to address the problem of overfitting³, we consider the function class \mathcal{F} (usually a class chosen with prior knowledge regarding the prediction problem) and restrict our predictors to come from this class, i.e. $\hat{f}_n \in \mathcal{F}$. We can then decompose our expected error as

$$L(\hat{f}_n) - L(f^B) = \underbrace{L(\hat{f}_n) - \inf_{f \in \mathcal{F}} L(f)}_{\text{estimation error (variance)}} + \underbrace{\inf_{f \in \mathcal{F}} L(f) - L(f^B)}_{\text{approximation error (bias)}}.$$

Here, \hat{f}_n shall be the predictor that minimizes the empirical error in our hypothesis class, i.e. $L_n(\hat{f}_n) = \inf_{f \in \mathcal{F}} L_n(f)$. In the sequel, we will only study the estimation error, i.e. we aim at statements of the form

$$\mathbb{P}\{L(\hat{f}_n) - L(f^*) > \epsilon\} \leq B(\epsilon, n, \mathcal{F}),$$

where B decreases to 0 as $n \rightarrow \infty$ for any ϵ and $f^* = \arg \inf_{f \in \mathcal{F}} L(f)$, i.e. the best predictor in the class \mathcal{F} (we will not study the approximation error here, details can for instance be found in [4]). Equivalently, we will consider the following kind of deviation inequality: With probability $1 - \delta$ for $\delta \in (0, 1)$ we want

$$L(\hat{f}_n) - L(f^*) \leq D(\delta, n, \mathcal{F}).$$

We will particularly be interested in statements for all $f \in \mathcal{F}$, i.e. we want uniform and not only pointwise bounds. Note that we have $L(\hat{f}_n) - L_n(\hat{f}_n) \leq \sup_{f \in \mathcal{F}} \{L(f) - L_n(f)\}$, so we will then also investigate the difference between expected and empirical loss, $L(\hat{f}_n) - L_n(\hat{f}_n)$, which we call generalization error⁴.

More abstractly, we consider the following learnability definition.

Definition 1.10 (PAC learnability [8]). *A hypothesis class \mathcal{F} is “probably approximately correctly” (PAC) learnable if there is an algorithm \mathcal{A} such that for all $\epsilon > 0$ and $\delta > 0$ there is a sample size $n(\epsilon, \delta) \in \mathbb{N}$ such that for all distributions \mathbb{P} over $\mathcal{X} \times \mathcal{Y}$*

$$\mathbb{P}(L(\mathcal{A}(S_n)) - L(f^*) < \epsilon) \geq 1 - \delta.$$

Remark 1.11. *One interesting aspect in introducing hypothesis classes \mathcal{F} is model selection. Consider a family of hypothesis classes $\{\mathcal{F}_\alpha\}_{\alpha \in A}$ with $\mathcal{F} = \cup_{\alpha \in A} \mathcal{F}_\alpha$ (e.g. $\mathcal{F}_1 \subset \mathcal{F}_2 \subset \dots$). Let \hat{f}_n^α be the predictor corresponding to the ERM over \mathcal{F}_α . One considers $\alpha_n := \alpha(S_n)$, i.e. a data-dependent selection of the possible classes. We are then interested in $L(\hat{f}_n^{\alpha_n}) - \inf_{\alpha \in A} \inf_{f \in \mathcal{F}_\alpha} L(f)$.*

Remark 1.12. *If we consider only one possible function, i.e. $\mathcal{F} = \{f_1\}$, then nothing will be learnt and we know from the law of large numbers that*

$$L_n(\hat{f}_n) - L(\hat{f}_n) = \frac{1}{n} \sum_{i=1}^n l(f_1(X_i), Y_i) - \mathbb{E}[l(f_1(X), Y)] \rightarrow 0$$

for $n \rightarrow \infty$ at rate $\frac{1}{\sqrt{n}}$.

So in the rather pathological case of $\mathcal{F} = \{f_1\}$ we have convergence of $L(\hat{f}_n) - L(f^*)$. But does it also hold with N predictors at hand, and if yes, how fast does it happen? To answer these questions the following so-called concentration inequalities will be of help.

1.1 Concentration inequalities

Concentration inequalities are statements of the form⁵

$$\mathbb{P}(|\bar{Z}_n - \mathbb{E}[Z]| < \epsilon) \geq 1 - \delta.$$

Theorem 1.13 (Markov inequality). *Consider a random variable Z with $Z \geq 0$. Then for all $\epsilon > 0$ we have*

$$\mathbb{P}(Z > \epsilon) \leq \frac{\mathbb{E}[Z]}{\epsilon}.$$

³Another way to address overfitting is to modify the criterion to be minimized by for instance adding a penalty term for too “complicated” functions – see structural risk minimization and (normalized) regularization.

⁴Compare to the uniform law of large numbers.

⁵Sometimes the notation $Pf = \mathbb{E}[f(X, Y)]$ and $P_n f = \frac{1}{n} \sum_{i=1}^n f(X_i, Y_i)$ for the empirical measure is used. Compare this to the theory of empirical processes, where uniform deviations of averages from their expectations are investigated, i.e. $\sup_{f \in \mathcal{F}} \{Pf - P_n f\}$.

Proof. See [6]. □

Theorem 1.14 (Chebysheff inequality). *Consider a random variable Z . Then for all $\epsilon > 0$ we have*

$$\mathbb{P}(|Z - \mathbb{E}[Z]| > \epsilon) \leq \frac{\text{Var}(Z)}{\epsilon^2}.$$

Proof. See [6]. □

We motivate the next bound with

Example 1.15. *Consider $X \sim \mathcal{N}(\mu, \sigma^2)$, then we have*

$$\mathbb{P}(|X - \mu| > \epsilon) = \int_{\epsilon}^{\infty} \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{x^2}{2\sigma^2}} dx \leq \int_{\epsilon}^{\infty} \frac{x}{\epsilon\sqrt{2\pi\sigma^2}} e^{-\frac{x^2}{2\sigma^2}} dx$$

and therefore $\mathbb{P}(|X - \mu| \geq \epsilon) \leq \frac{\sigma}{\epsilon} \sqrt{\frac{2}{\pi}} e^{-\frac{\epsilon^2}{2\sigma^2}}$. We will see that this is a special case of the following theorem, when replacing σ^2 with the maximal possible variance of a random variable in $[0, 1]$ (i.e. $\frac{1}{4}$). It is also intuitive to look at $\bar{X}_n \sim \mathcal{N}(\mu, \frac{\sigma^2}{n})$, i.e. $\mathbb{P}(|\bar{X}_n - \mu| > \epsilon) = \mathbb{P}(|Z| > \sqrt{n}\epsilon) \leq \frac{\sigma}{\sqrt{n}\epsilon} \sqrt{\frac{2}{\pi}} e^{-\frac{n\epsilon^2}{2\sigma^2}}$, which decreases exponentially for large n .

Lemma 1.16 (Hoeffding's lemma). *For a random variable $Z \in [a, b]$ with $\mathbb{E}[Z] = 0$ we can bound the moment generating function for all $\lambda \in \mathbb{R}$ as*

$$\mathbb{E}[e^{\lambda Z}] \leq \exp\left(\frac{\lambda^2(b-a)^2}{8}\right).$$

Proof. Since $e^{\lambda Z}$ is convex we have for all $a \leq Z \leq b$

$$e^{\lambda Z} \leq \frac{b-Z}{b-a} e^{\lambda a} + \frac{Z-a}{b-a} e^{\lambda b}$$

and therefore

$$\begin{aligned} \mathbb{E}[e^{\lambda Z}] &\leq \frac{b}{b-a} e^{\lambda a} - \frac{a}{b-a} e^{\lambda b} \\ &= \left(-\frac{a}{b-a}\right) e^{\lambda a} \left(-\frac{b}{a} + e^{\lambda b - \lambda a}\right) \\ &= \left(-\frac{a}{b-a}\right) e^{\lambda a} \left(-\frac{b-a}{a} - 1 + e^{\lambda(b-a)}\right) \\ &= (1 - c + ce^d) e^{-cd} \\ &=: e^{L(d)} \end{aligned}$$

with $c = -\frac{a}{b-a} > 0$, $d = \lambda(b-a)$ and

$$L(d) = -cd + \log(1 - c + ce^d).$$

Taking derivatives w.r.t. d brings

$$L(0) = L'(0) = 0$$

and

$$L''(d) = \frac{ce^d(1 - c + ce^d) - c^2 e^{2d}}{(1 - c + ce^d)^2} = \frac{ce^d}{1 - c + ce^d} \left(1 - \frac{ce^d}{1 - c + ce^d}\right) = t(1-t) \leq \frac{1}{4}$$

with $t = \frac{ce^d}{1 - c + ce^d} > 0$. By Taylor's theorem there exists an $e \in [0, d]$ s.t.

$$L(d) = L(0) + dL'(0) + \frac{1}{2}d^2L''(e) \leq \frac{1}{8}d^2 = \frac{\lambda^2}{8}(b-a)^2.$$

This implies the bound (by convexity of the logarithm). □

Theorem 1.17 (Hoeffding's inequality). *Let Z_1, \dots, Z_n be independent (but not necessarily identically distributed). Assume w.l.o.g. that $Z_i \in [0, 1]$ for $i = 1, \dots, n$. Then with $\bar{Z}_n = \frac{1}{n} \sum_{i=1}^n Z_i$*

$$\mathbb{P}(|\bar{Z}_n - \mathbb{E}[\bar{Z}_n]| > \epsilon) \leq 2e^{-2n\epsilon^2}.$$

Proof. Consider $Z_i \in [a_i, b_i]$. With Markov's inequality we have for $\lambda > 0$

$$\begin{aligned} \mathbb{P}(\bar{Z}_n - \mathbb{E}[\bar{Z}_n] > \epsilon) &\leq \frac{\mathbb{E}[e^{\lambda(\bar{Z}_n - \mathbb{E}[\bar{Z}_n])}]}{e^{\lambda\epsilon}} \\ &= \frac{\mathbb{E}[e^{\frac{\lambda}{n} \sum_{i=1}^n (Z_i - \mathbb{E}[Z_i])}]}{e^{\lambda\epsilon}} \\ &= \frac{\prod_{i=1}^n \mathbb{E}[e^{\frac{\lambda}{n} (Z_i - \mathbb{E}[Z_i])}]}{e^{\lambda\epsilon}} \\ &\leq \frac{\prod_{i=1}^n e^{\frac{\lambda^2}{8n^2} (b_i - a_i)^2}}{e^{\lambda\epsilon}} \\ &= \exp\left(\frac{\sum_{i=1}^n (b_i - a_i)^2}{8n^2} \lambda^2 - \lambda\epsilon\right), \end{aligned}$$

where in the last inequality we used Hoeffding's lemma. This general proof technique is called Chernoff bound. We now minimize w.r.t λ and get $\lambda^* = \frac{4n^2\epsilon}{\sum_{i=1}^n (b_i - a_i)^2}$ and therefore

$$\mathbb{P}(\bar{Z}_n - \mathbb{E}[\bar{Z}_n] > \epsilon) \leq \exp\left(-\frac{2n^2\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right),$$

which yields the desired result with $a_i = 0, b_i = 1$ and by noting that $\mathbb{P}(|Z| \leq \epsilon) \leq \mathbb{P}(Z \leq \epsilon) + \mathbb{P}(-Z \leq \epsilon)$. \square

Remark 1.18. *Note that this theorem also holds for functions of random variables. One can for instance consider $Z = g(X, Y)$ with $g(X, Y) = l(f(X), Y)$. However, also note that it only holds for a fixed function and not uniformly for all $f \in \mathcal{F}$. For a fixed sample one finds an $f \in \mathcal{F}$ that yields a very large error, as for instance seen in example 1.7.*

Remark 1.19. *We can formulate Hoeffding's bound as a deviation inequality by realizing the equivalence of the statements $\mathbb{P}(|Z| > \epsilon) \leq \delta$ and $\mathbb{P}(|Z| < \epsilon) \geq 1 - \delta$. Set $\delta := 2e^{-2n\epsilon^2}$, then $\epsilon = \sqrt{\frac{\log \frac{2}{\delta}}{2n}}$, and we have with probability at least $1 - \delta$ that*

$$|\bar{Z}_n - \mathbb{E}[\bar{Z}_n]| \leq \sqrt{\frac{\log \frac{2}{\delta}}{2n}}.$$

Hoeffding's inequality does not use any knowledge about the distribution of variables. Therefore one can consider Bernstein's inequality, which uses the variance of the distribution to get a tighter bound.

Theorem 1.20 (Bernstein's inequality). *Let Z_1, \dots, Z_n be independent (but not necessarily identically distributed). Assume w.l.o.g. $|Z_i| < 1$. Then with probability $1 - \delta$ we have*

$$|\bar{Z}_n - \mathbb{E}[\bar{Z}_n]| \leq \sqrt{2 \text{Var}(\bar{Z}_n) \log \frac{2}{\delta}} + \frac{2 \log \frac{2}{\delta}}{3n}.$$

Proof. See [2]. \square

Remark 1.21. *If the variance is very small then the first term on the right hand side becomes negligible.*

Theorem 1.22 (McDiarmid's inequality). *Consider a function $g : \mathcal{Z}^n \rightarrow \mathbb{R}$. Assume that there exist constants c_1, \dots, c_n s.t.*

$$\sup_{Z_1, \dots, Z_n, Z'} |g(Z_1, \dots, Z_n) - g(Z_1, \dots, Z_{i-1}, Z', Z_{i+1}, \dots, Z_n)| \leq c_i$$

for all $i = 1, \dots, n$. Let $Z_1, \dots, Z_n \sim \mathbb{P}$ i.i.d. with values in \mathcal{Z} and let $g := g(Z_1, \dots, Z_n)$. Then with probability $\geq 1 - \delta$ we have

$$|g - \mathbb{E}[g]| \leq \sqrt{\frac{1}{2} \left(\sum_{i=1}^n c_i^2 \right) \log \frac{2}{\delta}}.$$

Proof. First note that

$$\mathbb{P}(|g - \mathbb{E}[g]| \geq \epsilon) = \mathbb{P}(g - \mathbb{E}[g] \geq \epsilon) + \mathbb{P}(g - \mathbb{E}[g] \leq -\epsilon),$$

from which we will show the first inequality – the second one works analogously. Define for $i = 2, \dots, n$ the martingale difference sequence

$$V_i := \mathbb{E}[g|Z_1, \dots, Z_i] - \mathbb{E}[g|Z_1, \dots, Z_{i-1}] \quad \text{and} \quad V_1 := \mathbb{E}[g|Z_1] - \mathbb{E}[g],$$

then we have

$$g(Z_1, \dots, Z_n) - \mathbb{E}[g(Z_1, \dots, Z_n)] = \sum_{i=1}^n V_i.$$

With a similar argument as in Lemma 1.16 (see [2], section 6.1) we get

$$\mathbb{E}[e^{\lambda V_i} | Z_1, \dots, Z_{i-1}] \leq \exp\left(\frac{\lambda^2 c_i^2}{8}\right).$$

Then we have for all $\lambda > 0$

$$\begin{aligned} \mathbb{P}(g(Z_1, \dots, Z_n) - \mathbb{E}[g(Z_1, \dots, Z_n)] \geq \epsilon) &= \mathbb{P}\left(\sum_{i=1}^n V_i \geq \epsilon\right) \\ &= \mathbb{P}\left(e^{\lambda \sum_{i=1}^n V_i} \geq e^{\lambda \epsilon}\right) \\ &\leq e^{-\lambda \epsilon} \mathbb{E}\left[e^{\lambda \sum_{i=1}^n V_i}\right] \\ &= e^{-\lambda \epsilon} \mathbb{E}\left[e^{\lambda \sum_{i=1}^{n-1} V_i} \mathbb{E}[e^{\lambda V_n} | Z_1, \dots, Z_{n-1}]\right] \\ &\leq e^{-\lambda \epsilon} e^{\frac{\lambda^2 c_n^2}{8}} \mathbb{E}\left[e^{\lambda \sum_{i=1}^{n-1} V_i}\right] \\ &\leq e^{-\lambda \epsilon} e^{\frac{\lambda^2}{8} \sum_{i=1}^n c_i^2}. \end{aligned}$$

Taking $\lambda = \frac{4\epsilon}{\sum_{i=1}^n c_i^2}$ brings the statement. \square

Remark 1.23. Note that with $g(Z_1, \dots, Z_n) = \frac{1}{n} \sum_{i=1}^n Z_i$ we rediscover Hoeffdings inequality.

1.2 Error bounds for finite classes

Let us first consider the finite hypothesis space $\mathcal{F} = \{f_1, \dots, f_N\}$.

Theorem 1.24 (Realizable case). *Consider $\mathcal{Y} = \{0, 1\}$ in the zero-noise case, i.e. $L(f^B) = 0$ with $f^B \in \mathcal{F}$ (as in example 1.4). As before we call such an f^B , since it is in \mathcal{F} , just f^* . Then with probability $\geq 1 - \delta$ we have*

$$L(\hat{f}_n) - L(f^*) = L(\hat{f}_n) \leq \frac{\log N + \log\left(\frac{1}{\delta}\right)}{n}.$$

Proof. Let us try to bound $\mathbb{P}(\text{ERM selects } f \text{ with } L(f) > \epsilon)$. By definition we know that $L_n(\hat{f}_n) = 0$, since ERM selects only these predictors f for which $L_n(f) = 0$. Now take any $f \in \mathcal{F}$ s.t. $L(f) = \mathbb{P}(f(X) \neq Y) > \epsilon$, i.e. consider the functions that are expected to make some error. That implies that $\mathbb{P}(L(f) = 0) < 1 - \epsilon$ and therefore for n samples we get

$$\mathbb{P}(L_n(f) = 0) < (1 - \epsilon)^n \leq e^{-n\epsilon},$$

where the last step follows from $1 + x \leq e^x$. Now we consider

$$\begin{aligned} \mathbb{P}(\text{ERM selects } \hat{f}_n \text{ with } L(\hat{f}_n) > \epsilon) &\leq \mathbb{P}(\exists f \in \mathcal{F} \text{ s.t. } L(f) > \epsilon, L_n(f) = 0) \\ &= \mathbb{P}\left(\bigcup_{\substack{f \in \mathcal{F} \\ L(f) > \epsilon}} L_n(f) = 0\right) \\ &\leq \sum_{\substack{f \in \mathcal{F} \\ L(f) > \epsilon}} \mathbb{P}(L_n(f) = 0) \\ &\leq \sum_{\substack{f \in \mathcal{F} \\ L(f) > \epsilon}} e^{-n\epsilon} \\ &\leq N e^{-n\epsilon} =: \delta. \end{aligned}$$

This gives $\epsilon = \frac{\log N + \log(\frac{1}{\delta})}{n}$. So with probability at least $1 - \delta$ we have the complement event, i.e.

$$\mathbb{P}(\text{ERM selects } f \text{ with } L(f) \leq \epsilon) \geq 1 - \delta.$$

The event $\{\text{ERM selects } f \text{ with } L(f) \leq \epsilon\}$ is just $\{L(\hat{f}_n) \leq \epsilon\}$ and therefore we get the desired bound. \square

Remark 1.25. *One can show that given our assumptions this bound is actually the best possible one.*

Let us now drop the assumption $L(f^*) = 0$.

Theorem 1.26 (Agnostic case). *We consider again $\mathcal{F} = \{f_1, \dots, f_N\}$. Let l be bounded. Then with probability $\geq 1 - \delta$ we have*

$$L(\hat{f}_n) - L(f^*) \leq \sqrt{2 \frac{\log(2N) + \log(\frac{1}{\delta})}{n}}.$$

Proof.

$$\begin{aligned} L(\hat{f}_n) - L(f^*) &= L(\hat{f}_n) - L_n(\hat{f}_n) + \underbrace{L_n(\hat{f}_n) - L_n(f^*)}_{\leq 0} + L_n(f^*) - L(f^*) \\ &\leq |L(\hat{f}_n) - L_n(\hat{f}_n)| + |L_n(f^*) - L(f^*)| \\ &\leq 2 \sup_{f \in \mathcal{F}} |L(f) - L_n(f)|. \end{aligned}$$

Since our \mathcal{F} is finite, we can actually consider the maximum and compute

$$\begin{aligned} \mathbb{P}\left(\max_{i=1, \dots, N} |L(f_i) - L_n(f_i)| > \epsilon\right) &= \mathbb{P}\left(\cup_{i=1}^N \{|L(f_i) - L_n(f_i)| > \epsilon\}\right) \\ &\leq \sum_{i=1}^N \mathbb{P}(|L(f_i) - L_n(f_i)| > \epsilon) \\ &\leq 2Ne^{-2n\epsilon^2} =: \delta, \end{aligned}$$

where the last step follows from Hoeffding's inequality, which also holds when taking functions of the random variables. If we solve for ϵ we get $\epsilon = \sqrt{\frac{\log \frac{2N}{\delta}}{2n}}$. Considering again the complement event (as in the proof of theorem 1.24) yields our desired statement. \square

We have now seen the two extreme cases: The rather unrealistic realizable setting ($L(f^*) = 0$) and the totally agnostic setting ($L(f^*) = \frac{1}{2}$). The first one has convergence rate $\mathcal{O}(\frac{1}{n})$, the latter one $\mathcal{O}(\frac{1}{\sqrt{n}})$. Let us now look at cases in between those extremes (“from slow to fast rates”).

Theorem 1.27. *We again consider $\mathcal{F} = \{f_1, \dots, f_N\}$ and our loss being bounded. With probability $\geq 1 - \delta$ we have*

$$L(\hat{f}_n) - L(f^*) \leq C \left(\sqrt{L(f^*) \frac{\log 4N + \log(\frac{1}{\delta})}{n}} + \frac{\log(\frac{4N}{\delta})}{n} \right),$$

where C is a constant.

Proof. First note that

$$\text{Var}(\mathbb{1}_{\{f(X) \neq Y\}}) = \mathbb{E}[\mathbb{1}^2] - \mathbb{E}[\mathbb{1}]^2 \leq \mathbb{E}[\mathbb{1}] = L(f).$$

Now take Bernstein's inequality with $\delta \rightarrow \frac{\delta}{N}$:

$$\mathbb{P}\left(\exists f \in \mathcal{F} : \underbrace{|L_n(f) - L(f)|}_{=: E(f)} \geq \sqrt{2L(f) \frac{\log \frac{2N}{\delta}}{n}} + \frac{2 \log \frac{2N}{\delta}}{3n}\right) = \mathcal{P}(\cup_{f \in \mathcal{F}} E(f)) \leq \sum_{f \in \mathcal{F}} \mathbb{P}(E) \leq \sum_{f \in \mathcal{F}} \frac{\delta}{N} = \delta.$$

Forming the complement yields

$$\mathbb{P} \left(\forall f \in \mathcal{F} : |L_n(f) - L(f)| \leq \sqrt{2L(f) \frac{\log \frac{2N}{\delta}}{n}} + \frac{2 \log \frac{2N}{\delta}}{3n} \right) \geq 1 - \delta.$$

We can therefore pick $f = \hat{f}_n$ and get that with probability at least $1 - \delta$

$$|L_n(\hat{f}_n) - L(\hat{f}_n)| \leq \sqrt{2L(\hat{f}_n) \frac{\log \frac{2N}{\delta}}{n}} + \frac{2 \log \frac{2N}{\delta}}{3n}.$$

Using another time Bernstein's inequality for f^* yields

$$L_n(\hat{f}_n) \leq L_n(f^*) \leq L(f^*) + \sqrt{2L(f^*) \frac{\log \frac{2}{\delta}}{n}} + \frac{2 \log \frac{2}{\delta}}{3n}$$

We can now combine the preceding two equations to get

$$\frac{L(\hat{f}_n) - L(f^*)}{\sqrt{L(\hat{f}_n)}} \leq \sqrt{2 \frac{\log \frac{4N}{\delta}}{n}}$$

and use the fact $\frac{A-B}{\sqrt{A}} \leq C \Rightarrow A \leq B + C^2 + \sqrt{BC}$ for $A, B, C \geq 0$ to get the result. \square

Remark 1.28. *This bound is again tight. Note that if we assume $n < \frac{1}{L(f^*)} =: \gamma$ we are in a fast-rates-regime, i.e. $\mathcal{O}(\frac{1}{n})$.*

Remark 1.29. *Note that we often do not use ERM in practice, as the minimization can be NP-hard.*

1.3 Error bounds for infinite classes

We now want to consider function classes \mathcal{F} that have infinitely many elements. The former bounds will then be meaningless (as we have $N = \infty$). When looking into the proofs, we realize that our union bounds as in theorem 1.26 will yield an infinite sum that has to be bounded in order to be useful. If the function class is even uncountable we need a completely different approach. For that, we make the following definition.

Definition 1.30 (Rademacher complexity). *Consider a function class \mathcal{F} and i.i.d. random variables $Z_1, \dots, Z_n \sim \mathbb{P}$. Let $\sigma = \{\sigma_1, \dots, \sigma_n\}$ be a set of i.i.d. random variables s.t. $\mathbb{P}(\sigma_i = \pm 1) = \frac{1}{2}$. We then define the empirical Rademacher complexity⁶ as*

$$\hat{R}_n(\mathcal{F}, S_n) = \mathbb{E}_\sigma \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \right| \middle| S_n \right].$$

For the Rademacher complexity we average over all samples $S_n = (Z_1, \dots, Z_n)$, i.e.

$$R_n(\mathcal{F}) = \mathbb{E}_{S_n} [\hat{R}_n(\mathcal{F})].$$

Remark 1.31. *The supremum can be interpreted as finding the function f out of \mathcal{F} that makes the samples correlate the most with some random values σ . Loosely speaking, the idea is to find a function that looks like random noise the most. The Rademacher complexity then gives a measure of how well we can find a function out of the function class \mathcal{F} that behaves like random noise. Note that $\hat{R}_n(\mathcal{F}) = 0$ if \mathcal{F} consists of only one function and $\hat{R}_n(\mathcal{F}) = 1$ if \mathcal{F} consists of all functions. In fact, we have $\hat{R}_n(\mathcal{F}) \in [0, 1]$. Note that in general computing Rademacher complexities can be difficult (in fact, as difficult as computing empirical risk minimizers).*

Remark 1.32. *For $\mathcal{F} \subset \mathcal{G}$ we have $R_n(\mathcal{F}) \leq R_n(\mathcal{G})$.*

⁶In the context of empirical processes $R_n(f) = \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i)$ is sometimes called Rademacher process.

Theorem 1.33. Consider any prediction task $(\mathbb{P}, \mathcal{X}, \mathcal{Y}, l, \mathcal{F})$ s.t. $l(f(X), Y) \in [0, 1]$. Then with probability $\geq 1 - \delta$

$$\sup_{f \in \mathcal{F}} |L(f) - L_n(f)| \leq 2R_n(l \circ \mathcal{F}) + \sqrt{\frac{\log \frac{1}{\delta}}{2n}},$$

where $l \circ \mathcal{F} = \{l(f(X), Y) : f \in \mathcal{F}, (X, Y) \in \mathcal{X} \times \mathcal{Y}\}$.

Proof. Comparing to theorem 1.22, let $g(Z_1, \dots, Z_n) = \sup_{f \in \mathcal{F}} |L_n(f) - L(f)|$. If we change one datum Z_i to Z'_i we get

$$\begin{aligned} \left| \sup_{f \in \mathcal{F}} |L_n(f) - L(f)| - \sup_{f \in \mathcal{F}} |L_n^i(f) - L(f)| \right| &\leq \left| \frac{1}{n} \sup_{f \in \mathcal{F}} |l(f(X_i), Y_i) - l(f(X'_i), Y'_i)| \right| \\ &\leq \frac{1}{n}, \end{aligned}$$

where L_n^i is the empirical risk with the changed datum. We can therefore use McDiarmid's inequality

$$\sup_{f \in \mathcal{F}} |L_n(f) - L(f)| \leq \mathbb{E} \left[\sup_{f \in \mathcal{F}} |L_n(f) - L(f)| \right] + \sqrt{\frac{\log \frac{1}{\delta}}{2n}}.$$

The idea now is to introduce a ghost sample S'_n that is identically distributed as S_n . We then have

$$\begin{aligned} \mathbb{E}_{S_n} \left[\sup_{f \in \mathcal{F}} |L_n(f) - L(f)| \right] &= \mathbb{E}_{S_n} \left[\sup_{f \in \mathcal{F}} |\mathbb{E}_{S'_n} [L_n(f) - L'_n(f)]| \right] \\ &\leq \mathbb{E}_{S_n, S'_n} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n (l(f(X_i), Y_i) - l(f(X'_i), Y'_i)) \right| \right] \\ &= \mathbb{E}_{S_n, S'_n, \sigma} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i (l(f(X_i), Y_i) - l(f(X'_i), Y'_i)) \right| \right] \\ &\leq \mathbb{E}_{S_n, \sigma} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i l(f(X_i), Y_i) \right| \right] + \mathbb{E}_{S'_n, \sigma} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i l(f(X'_i), Y'_i) \right| \right] \\ &= 2R_n(l \circ \mathcal{F}), \end{aligned}$$

where we used Jensen's inequality in the second step. □

Remark 1.34. Note by looking at the proof of theorem 1.26 that we actually have

$$L(\hat{f}_n) - L(f^*) \leq \sup_{f \in \mathcal{F}} (L(f) - L_n(f)) + \sup_{f \in \mathcal{F}} (L_n(f) - L(f))$$

and can therefore bound the estimation error.

Let us now investigate how to control the Rademacher complexity. First, we rediscover the finite class case.

Lemma 1.35. Assume $\mathcal{F} = \{f_1, \dots, f_N\}$. Then

$$\hat{R}_n(\mathcal{F}) \leq \frac{\sqrt{2 \log N}}{n} \max_{j=1, \dots, N} \sqrt{\sum_{i=1}^n f_j^2(Z_i)}$$

and therefore

$$L(\hat{f}_n) - L(f^*) \leq C \left(\sqrt{\frac{\log N}{n}} + \sqrt{\frac{\log \frac{1}{\delta}}{n}} \right).$$

Proof. We have for any $\lambda > 0$

$$\begin{aligned}
\exp \left(\lambda \mathbb{E} \left[\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \sigma_i f_j(Z_i) \right] \right) &\leq \mathbb{E} \left[\exp \left(\lambda \max_{j=1, \dots, N} \frac{1}{n} \sum_{i=1}^n \sigma_i f_j(Z_i) \right) \right] \\
&= \mathbb{E} \left[\max_{j=1, \dots, N} \exp \left(\frac{\lambda}{n} \sum_{i=1}^n \sigma_i f_j(Z_i) \right) \right] \\
&\leq \sum_{j=1}^N \prod_{i=1}^n \mathbb{E} \left[\exp \left(\frac{\lambda}{n} \sigma_i f_j(Z_i) \right) \right] \\
&\leq \sum_{j=1}^N \prod_{i=1}^n \exp \left(\frac{\lambda^2}{2n^2} f_j^2(Z_i) \right) \\
&= \sum_{j=1}^N \exp \left(\frac{\lambda^2}{2n^2} \sum_{i=1}^n f_j^2(Z_i) \right) \\
&\leq N \max_{j=1, \dots, N} \exp \left(\frac{\lambda^2}{2n^2} \sum_{i=1}^n f_j^2(Z_i) \right),
\end{aligned}$$

where we used Hoeffding's lemma and which yields that

$$\hat{R}_n(\mathcal{F}) \leq \frac{1}{\lambda} \left(\log N + \max_{j=1, \dots, N} \frac{\lambda^2}{2n^2} \sum_{i=1}^n f_j^2(Z_i) \right).$$

Minimizing w.r.t. λ brings $\lambda^* = \sqrt{\frac{\log N 2n^2}{\max_{j=1, \dots, N} \sum_{i=1}^n f_j^2(Z_i)}}$ and therefore the result. \square

Remark 1.36. Sometimes the Rademacher complexity is defined without the absolute values. Note, however, that we have $\mathbb{E} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \right| \right] = \mathbb{E} \left[\sup_{f \in \mathcal{F} \cup -\mathcal{F}} \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \right]$.

Let us now consider the infinite class case, i.e. a general \mathcal{F} for binary classification with $l(y', y) = \mathbb{1}\{y' \neq y\}$. From theorem 1.33 we know that

$$L(\hat{f}_n) - L(f^*) \leq 4R_n(l \circ \mathcal{F}) + \sqrt{\frac{2 \log \frac{1}{\delta}}{n}}.$$

Still, in R_n we take the supremum over all $f \in \mathcal{F}$, thus we need to find a strategy how to deal with that in the infinite class case. The idea is to argue that we can divide \mathcal{F} into finitely many equivalence classes. Note that in the computation of R_n we consider the sum over n summands, where each of them has the value 0 or 1. We can say that each function f induces an n -dimensional vector with the losses of the sample,

$$f \mapsto (l(f(X_1), Y_1), \dots, l(f(X_n), Y_n)) \in \{0, 1\}^n,$$

and we realize that we can have at most 2^n such vectors, i.e. no matter how big the function class is, we have to consider at most 2^n cases. Therefore, we can divide \mathcal{F} into 2^n equivalence classes and identify functions that lead to the same loss vector with one another. Unfortunately, however, this is not enough, as plugging in 2^n instead of N in lemma 1.35 just yields a constant bound that is not decreasing with n . The next idea is to consider only function classes \mathcal{F} that yield (much) less than 2^n possible permutations. Let us therefore define the following set of vectors, the projections of the predictors on the sample,

$$\mathcal{F}_{S_n} := \{(l(f(X_1), Y_1), \dots, l(f(X_n), Y_n)), f \in \mathcal{F}\} \subset \{0, 1\}^n,$$

and ask what the cardinality of the vector over all possible data sets S_n is. For this we define the so called growth function (or shattering number)

$$S_{\mathcal{F}}(n) = \sup_{S_n} |\mathcal{F}_{S_n}| \leq 2^n,$$

which is the worst case cardinality of our equivalence classes. The hope is that the growth function will be indeed much smaller than 2^n . Note that $S_{\mathcal{F}}(n)$ does not depend on the sample distribution \mathbb{P} , but only on \mathcal{F} .

Now we introduce a crucial trick.

Lemma 1.37 (Symmetrization). *For all $\epsilon \geq \sqrt{\frac{2}{n}}$*

$$\mathbb{P} \left(\sup_{f \in \mathcal{F}} |L_n(f) - L(f)| > \epsilon \right) \leq 2\mathbb{P} \left(\sup_{f \in \mathcal{F}} |L_n(f) - L'_n(f)| > \frac{\epsilon}{2} \right),$$

where L'_n is the empirical error of a ghost sample.

Proof. We again consider a ghost sample $S'_n = (Z'_1, \dots, Z'_n)$ that is distributed as the original sample S_n and denote its empirical error as L'_n . Let $f_n \in \mathcal{F}$ maximize $|L_n(f) - L(f)|$. Note that if $|L_n(f_n) - L(f_n)| > \epsilon$ and $|L(f_n) - L'_n(f_n)| \leq \frac{\epsilon}{2}$, then

$$\epsilon < |L_n(f_n) - L(f_n)| = |L_n(f_n) - L'_n(f_n) + L'_n(f_n) - L(f_n)| \leq |L_n(f_n) - L'_n(f_n)| + |L'_n(f_n) - L(f_n)| \leq \frac{\epsilon}{2} + |L'_n(f_n) - L(f_n)|$$

and therefore $|L'_n(f_n) - L(f_n)| > \frac{\epsilon}{2}$. This then yields

$$\begin{aligned} \mathbb{1} \left\{ |L_n(f_n) - L(f_n)| > \epsilon \right\} \mathbb{1} \left\{ |L(f_n) - L'_n(f_n)| \leq \frac{\epsilon}{2} \right\} &= \mathbb{1} \left\{ |L_n(f_n) - L(f_n)| > \epsilon, |L(f_n) - L'_n(f_n)| \leq \frac{\epsilon}{2} \right\} \\ &\leq \mathbb{1} \left\{ |L(f_n) - L'_n(f_n)| > \frac{\epsilon}{2} \right\}. \end{aligned}$$

Taking the expectation w.r.t. S'_n brings

$$\mathbb{1} \left\{ |L_n(f_n) - L(f_n)| > \epsilon \right\} \mathbb{P}' \left(|L(f_n) - L'_n(f_n)| \leq \frac{\epsilon}{2} \right) \leq \mathbb{P}' \left(|L_n(f_n) - L'_n(f_n)| > \frac{\epsilon}{2} \right).$$

Note that by Chebyshev's inequality and with $X \in [0, 1] \Rightarrow \text{Var}(X) \leq \frac{1}{4}$ we get

$$\mathbb{P}' \left(|L(f_n) - L'_n(f_n)| \leq \frac{\epsilon}{2} \right) \leq \frac{4 \text{Var}'(f_n)}{n\epsilon^2} \leq \frac{1}{n\epsilon^2} \leq \frac{1}{2},$$

which then yields

$$\mathbb{1} \left\{ \sup_{f \in \mathcal{F}} |L_n(f) - L(f)| > \epsilon \right\} \leq 2\mathbb{P}' \left(\sup_{f \in \mathcal{F}} |L_n(f) - L'_n(f)| > \frac{\epsilon}{2} \right).$$

Taking expectation w.r.t. S_n brings the desired result. \square

Remark 1.38. *Note that $L_n(f) - L(f)$ can take any real value, whereas $L_n(f) - L'_n(f)$ can only take finitely many values, i.e. we can hope to apply the union bound in the uncountable case.*

Theorem 1.39 (Vapnik-Chervonenkis bound). *We again consider $\mathcal{Y} = \{0, 1\}$. Then with probability $\geq 1 - \delta$*

$$L(\hat{f}_n) - L(f^*) \leq 2\sqrt{\frac{8}{n} \log \left(\frac{4S_{\mathcal{F}}(2n)}{\delta} \right)}.$$

Proof. With the preceding lemma we have

$$\begin{aligned} \mathbb{P} \left(\sup_{f \in \mathcal{F}} |L_n(f) - L(f)| \geq \epsilon \right) &\leq 2\mathbb{P} \left(\sup_{f \in \mathcal{F}} |L_n(f) - L'_n(f)| > \frac{\epsilon}{2} \right) \\ &= 2\mathbb{P} \left(\max_{l \in \mathcal{F}_{S_n, S'_n}} |L_n(l) - L'_n(l)| > \frac{\epsilon}{2} \right) \\ &\leq 2 \sum_{l \in \mathcal{F}_{S_n, S'_n}} \mathbb{P} \left(|L_n(l) - L'_n(f)| > \frac{\epsilon}{2} \right) \\ &\leq 2 \sum_{l \in \mathcal{F}_{S_n, S'_n}} 2e^{-\frac{n}{8}\epsilon^2} \\ &\leq 4S_{\mathcal{F}}(2n)e^{-\frac{n}{8}\epsilon^2}, \end{aligned}$$

where with some abuse of notation $L_n(l) - L'_n(l) = \frac{1}{n} \left(\sum_{i=1}^n l_i - \sum_{i=n+1}^{2n} l_i \right)$ with $l_i = l(f(X_i), Y_i)$, for which we used Hoeffding's inequality. \square

Still we have to answer the question when this bound goes to zero for $n \rightarrow \infty$, or in other words, what the geometry of \mathcal{F} has to be s.t. $S_{\mathcal{F}}(n)$ does not grow too fast. The crucial tool that shall answer these questions is the following.

Definition 1.40 (VC dimension). *The VC dimension of a function class \mathcal{F} is defined as*

$$VC(\mathcal{F}) = \sup \{n : S_{\mathcal{F}}(n) = 2^n\}.$$

If such an n does not exist we set $VC(\mathcal{F}) = \infty$.

Remark 1.41. $VC(\mathcal{F}) = d$ implies that there exists a sample $X_1, \dots, X_d \in \mathcal{X}$ s.t. for any $Y_1, \dots, Y_d \in \mathcal{Y}$ there exists an $f \in \mathcal{F}$ with $f(X_i) = Y_i$. We say that the sample is “shattered”.

Remark 1.42. The VC dimension can be understood as measuring the effective size of a function class. Note that it is distribution independent. One can consider a concept called VC entropy to include properties of \mathbb{P} (see for instance chapter 4.3 in [3]).

Class \mathcal{F}	$VC(\mathcal{F})$
$\{f_1, \dots, f_N\}$	$\leq \log_2 N$
Intervals on the real line	2
Discs in \mathbb{R}^2	3
Rectangles in \mathbb{R}^d	$2d$
Half-spaces in \mathbb{R}^d	$d + 1$
Convex polygons in \mathbb{R}^2	∞
$\{\text{sgn}(\sin(\alpha x)) : \alpha \in \mathbb{R}\}$	∞

Figure 1: Examples of VC dimensions.

Example 1.43 (VC dimensions). *The last example in figure 1 is interesting as it shows that the VC dimension is not necessarily linked to the number of free parameters. To show this take for instance $X_i = 10^{-i}$ and $\alpha = \pi \left(1 + \sum_{i=1}^n \frac{1}{2}(1 - Y_i)10^i\right)$. Then*

$$\begin{aligned}
f(X_j) &= \text{sgn} \left(\sin \left(10^{-j} \pi + \sum_{i=1}^n \frac{\pi}{2} (1 - Y_i) 10^{i-j} \right) \right) \\
&= \text{sgn} \left(\sin \left(10^{-j} \pi + \sum_{i \leq j, Y_i = -1} \frac{\pi}{2} (1 - Y_i) 10^{i-j} \right) \right) \\
&= \text{sgn} \left(\sin \left(\frac{\pi}{2} (1 - Y_j) + \pi \underbrace{\left(10^{-j} + \sum_{i < j, Y_i = -1} 10^{i-j} \right)}_{< 1} \right) \right),
\end{aligned}$$

after the properties of the geometric series. Now, if $Y_j = 1$ the sine function takes values between 0 and π and if $Y_j = -1$ it takes values between π and 2π – therefore the classification is correct.

The VC dimension brings the following theorem.

Theorem 1.44 (Sauer, Shelah, Perles, Vapnik, Chervonenkis). *Assume $VC(\mathcal{F}) = d < \infty$. Then for all n*

$$S_{\mathcal{F}}(n) \leq \sum_{i=0}^d \binom{n}{i}$$

and for all $n \geq d$ we get

$$S_{\mathcal{F}}(n) \leq \left(\frac{en}{d} \right)^d.$$

Proof. For $n = d = 1$ we have

$$S_{\mathcal{F}}(1) \leq \binom{1}{0} + \binom{1}{1} = 2,$$

which clearly holds. By induction assume that the statement holds for $n - 1$ and $d - 1$ as well as for $n - 1$ and d . We want to show that it also holds for n and d . Let $S_1 = (Z_1, \dots, Z_n)$ and $S_2 = (Z_2, \dots, Z_n)$, let \mathcal{F}_{S_1} and \mathcal{F}_{S_2} be the corresponding projections onto \mathcal{F} . For $f, g \in \mathcal{F}$ write

$$f \sim g \quad \text{if} \quad f(Z_1) = 1 - g(Z_1) \quad \text{and} \quad f(Z_j) = g(Z_j) \quad \text{for } j = 2, \dots, n.$$

and let

$$\mathcal{G} = \{f \in \mathcal{F} : \exists g \in \mathcal{F} \text{ s.t. } f \sim g\}.$$

Now consider \mathcal{G}_{S_2} , then

$$|\mathcal{F}_{S_1}| = |\mathcal{F}_{S_2}| + |\mathcal{G}_{S_2}|.$$

Note that $\text{VC}(\mathcal{F}_{S_2}) \leq d$ and $\text{VC}(\mathcal{G}_{S_2}) \leq d - 1$. The latter follows since, if \mathcal{G}_{S_2} shatters a set, we can add Z_1 to create a set that is shattered by \mathcal{F}_{S_1} . Therefore we have by assumption that

$$\begin{aligned} |\mathcal{F}_{S_1}| &\leq \sum_{i=0}^d \binom{n-1}{i} + \sum_{i=0}^{d-1} \binom{n-1}{i} \\ &= \sum_{i=0}^d \left(\binom{n-1}{i} + \binom{n-1}{i-1} \right) \\ &= \sum_{i=0}^d \binom{n}{i}. \end{aligned}$$

To prove the second statement we note that for $n \geq d$

$$\begin{aligned} \sum_{i=0}^d \binom{n}{i} &\leq \left(\frac{n}{d}\right)^d \sum_{i=0}^d \binom{n}{i} \left(\frac{d}{n}\right)^i \\ &\leq \left(\frac{n}{d}\right)^d \sum_{i=0}^n \binom{n}{i} \left(\frac{d}{n}\right)^i \\ &= \left(\frac{n}{d}\right)^d \left(1 + \frac{d}{n}\right)^n \\ &\leq \left(\frac{n}{d}\right)^d e^d. \end{aligned}$$

□

In analogy to the Vapnik-Chervonenkis bound one can also bound the Rademacher complexity. For this we need

Lemma 1.45. *Let X_1, \dots, X_n be random variables with $\mathbb{E}[e^{\lambda X_i}] \leq e^{\frac{\lambda^2 \xi^2}{2}}$ for all $\lambda > 0$ and $\xi > 0$. Then*

$$\mathbb{E} \left[\max_{1 \leq i \leq n} X_i \right] \leq \xi \sqrt{2 \log n}.$$

Proof. With Jensen's inequality we have

$$\begin{aligned} \exp \left(\lambda \mathbb{E} \left[\max_{1 \leq i \leq n} X_i \right] \right) &\leq \mathbb{E} \left[\exp \left(\lambda \max_{1 \leq i \leq n} X_i \right) \right] \\ &= \mathbb{E} \left[\max_{1 \leq i \leq n} \exp(\lambda X_i) \right] \\ &\leq \sum_{i=1}^n \mathbb{E} [e^{\lambda X_i}] \\ &\leq n e^{\frac{\lambda^2 \xi^2}{2}}. \end{aligned}$$

Therefore,

$$\mathbb{E} \left[\max_{1 \leq i \leq n} X_i \right] \leq \frac{\log n}{\lambda} + \frac{\lambda \xi^2}{2}$$

and minimizing w.r.t. λ gives the desired bound. \square

Theorem 1.46. *Let \mathcal{F} be a set of binary functions. Then*

$$R_n(\mathcal{F}) \leq \sqrt{\frac{2 \log S_{\mathcal{F}}(n)}{n}}.$$

Proof. We have

$$\begin{aligned} R_n(\mathcal{F}) &= \mathbb{E} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \right| \right] \\ &= \mathbb{E}_{S_n} \left[\mathbb{E}_{\sigma} \left[\sup_{f \in \mathcal{F}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \right| \middle| S_n \right] \right] \\ &= \mathbb{E} \left[\mathbb{E} \left[\max_{v \in \mathcal{F}_{S_n}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i v_i \right| \middle| S_n \right] \right]. \end{aligned}$$

Note that $\frac{\sigma v_i}{n}$ has mean zero and $-\frac{1}{n} \leq \frac{\sigma v_i}{n} \leq \frac{1}{n}$. Therefore, with lemma 1.16 we have $\mathbb{E} [e^{\lambda \sigma_i v_i}] \leq e^{\frac{\lambda^2}{2n^2}}$ and $\mathbb{E} [e^{\lambda \sum_{i=1}^n \frac{\sigma v_i}{n}}] \leq e^{\frac{\lambda^2}{2n}}$ and thus with the preceding lemma 1.45 we get

$$\mathbb{E} \left[\max_{v \in \mathcal{F}_{S_n}} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i v_i \right| \middle| S_n \right] \leq \sqrt{\frac{2 \log |\mathcal{F}_{S_n}|}{n}} \leq \sqrt{\frac{2 \log S_{\mathcal{F}}(n)}{n}}.$$

\square

Remark 1.47. *For $n \geq d$ the growth function $S_{\mathcal{F}}(n)$ only grows polynomially and therefore gives a helpful bound when plugging it in to theorem 1.39, namely $\leq C \sqrt{\frac{d \log n}{n}}$. In other words, a finite VC dimension brings learnability. In fact, using techniques called “chaining” and “covering numbers” one can even get a bound $\leq C \sqrt{\frac{d}{n}}$ – details can for instance be found in [5], chapter 4.3.*

In analogy to theorem 1.27 one can show

Theorem 1.48. *With probability at least $1 - \delta$ we have*

$$L(\hat{f}_n) - L(f^*) \leq 2 \sqrt{L(f^*) \frac{\log S_{\mathcal{F}}(2n) + \log \frac{4}{\delta}}{n}} + \frac{\log S_{\mathcal{F}}(2n) + \log \frac{4}{\delta}}{n}.$$

Proof. See chapter 6.2 in [3]. \square

Lemma 1.49 (Contraction principle). *Assume $\varphi : \mathbb{R} \rightarrow \mathbb{R}, \varphi(0) = 0$ and φ is L_{φ} -Lipschitz. Then for all function classes \mathcal{F}*

$$R_n(\varphi \circ \mathcal{F}) \leq L_{\varphi} R_n(\mathcal{F}).$$

Proof. Let $g_i(z)$ and $h_i(z)$ be functions s.t. for any z, z'

$$|g_i(z) - g_i(z')| \leq |h_i(z) - h_i(z')|.$$

Let us show

$$\mathbb{E}_{\sigma_1, \dots, \sigma_n, x} \left[\sup_z \left\{ c(x, z) + \sum_{i=1}^n \sigma_i g_i(z) \right\} \right] \leq \mathbb{E}_{\sigma_1, \dots, \sigma_n, x} \left[\sup_z \left\{ c(x, z) + \sum_{i=1}^n \sigma_i h_i(z) \right\} \right]$$

by induction. Obviously it is true for $n = 0$. Let us assume it holds for n , then

$$\begin{aligned}
\mathbb{E}_{\sigma_1, \dots, \sigma_{n+1}, x} \left[\sup_z \left\{ c(x, z) + \sum_{i=1}^{n+1} \sigma_i g_i(z) \right\} \right] &= \mathbb{E}_{\sigma_1, \dots, \sigma_n, x} \left[\sup_{z, z'} \left\{ \frac{c(x, z) + c(x, z')}{2} + \sum_{i=1}^n \sigma_i \frac{g_i(z) - g_i(z')}{2} + \frac{g_{n+1}(z) - g_{n+1}(z')}{2} \right\} \right] \\
&\leq \mathbb{E}_{\sigma_1, \dots, \sigma_n, x} \left[\sup_{z, z'} \left\{ \frac{c(x, z) + c(x, z')}{2} + \sum_{i=1}^n \sigma_i \frac{g_i(z) - g_i(z')}{2} + \frac{h_{n+1}(z) - h_{n+1}(z')}{2} \right\} \right] \\
&= \mathbb{E}_{\sigma_1, \dots, \sigma_{n+1}, x} \left[\sup_y \left\{ c(x, z) + \sum_{i=1}^n \sigma_i g_i(y) + \sigma_{n+1} h_{n+1}(z) \right\} \right] \\
&\leq \mathbb{E}_{\sigma_1, \dots, \sigma_{n+1}, x} \left[\sup_z \left\{ c(x, z) + \sum_{i=1}^{n+1} \sigma_i h_i(z) \right\} \right].
\end{aligned}$$

One can show the same also with absolute values. Now take $c(x, z) = 0$, $f(z) = \varphi(f(x)y)$ and $h(z) = L_\varphi f(x)y$ to get the result. \square

Remark 1.50. This lemma brings that Rademacher complexity is bounded whenever the loss is Lipschitz and therefore also brings learning statements for the regression setting. To give some examples, hinge loss given by $l(y, y') = \max\{1 - yy', 0\}$ is 1-Lipschitz, logistic loss given by $l(y, y') = \log(1 + e^{-yy'})$ is 1-Lipschitz, squared loss $l(y, y') = (y - y')^2$ is 4B-Lipschitz when $|y|, |y'| \leq B$, absolute loss $l(y, y') = |y - y'|$ is 1-Lipschitz.

- the concepts used in defining VC dimension apply only to binary classification, but we are often interested in generalization bounds for multiclass classification and regression as well
- For the regression setting a concept called “fat shattering” has been developed, however, we can use the Rademacher complexity directly in this case
- check <http://www.cs.cmu.edu/~ninamf/ML11/lect1117.pdf>
- more on the Rademacher complexity can be found in [1]
- check out On the Complexity of Linear Prediction: Risk Bounds, Margin Bounds, and Regularization
- check out <https://www.cs.cornell.edu/~sridharan/rad-paper.pdf>

2 Algorithms

2.1 Perceptron algorithm

Input: $(X_1, Y_1), \dots, (X_n, Y_n) \in \mathbb{R}^d \times \{-1, 1\}$
 $\theta_0 = (0, \dots, 0)^\top \in \mathbb{R}^d, t = 0$
while some (X_i, Y_i) is misclassified, i.e. $Y_i \neq \text{sgn}(\theta_t^\top X_i)$ **do**
 pick some misclassified (X_i, Y_i) at random
 $\theta_{t+1} \leftarrow \theta_t + Y_i X_i$
 $t \leftarrow t + 1$
end while

Theorem 2.1. Given that the data is linearly separable, i.e. there is a θ s.t. $Y_i \theta^\top X_i > 0$ for all i , the algorithm terminates (with zero empirical risk) after at most $\frac{R^2}{\gamma^2}$ iterations, where $R = \max_i \|X_i\|$ and $\gamma = \min_i \frac{\theta^\top X_i Y_i}{\|\theta\|}$,

Proof. We have

$$\theta_{t+1}^\top \theta = (\theta_t + Y_i X_i)^\top \theta \geq \theta_t^\top \theta + \gamma \|\theta\|$$

and since $\theta_0 = 0$ this brings $\theta_t^\top \theta \geq t\gamma \|\theta\|$. On the other hand, since X_i was incorrectly classified,

$$\|\theta_{t+1}\|^2 = \|\theta_t + Y_i X_i\|^2 = \|\theta_t\|^2 + 2Y_i \theta_t^\top X_i + \|X_i\|^2 \leq \|\theta_t\|^2 + R^2$$

so $\|\theta_t\|^2 \leq tR^2$. Combining those we get with Cauchy-Schwarz

$$t\gamma \|\theta\| \leq \theta_t^\top \theta \leq \|\theta_t\| \|\theta\| \leq \sqrt{t} R \|\theta\|$$

and therefore the statement. \square

Remark 2.2. As seen before the VC dimension of the perceptron in \mathbb{R}^d is $d + 1$.

Theorem 2.3. For a class $\mathcal{F} = \{\theta^\top x \in [-1, 1], x \in \mathbb{R}^n, \theta \in \mathbb{R}^n, \|\theta\|_2 \leq A, \|x\|_2 \leq B\}$ we have

$$\hat{R}_n(\mathcal{F}, S_n) \leq \frac{AB}{\sqrt{n}}$$

Proof.

$$\begin{aligned} \hat{R}_n(\mathcal{F}, S_n) &= \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{f \in \mathcal{F}} \left| \sum_{i=1}^n \sigma_i f(X_i) \right| \right] = \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{\|\theta\|_2 \leq A} \left| \sum_{i=1}^n \sigma_i \theta^\top X_i \right| \right] \\ &\leq \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{\|\theta\|_2 \leq A} \left| \theta^\top \sum_{i=1}^n \sigma_i X_i \right| \right] \leq \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{\|\theta\|_2 \leq A} \|\theta\| \left\| \sum_{i=1}^n \sigma_i X_i \right\| \right] \\ &= \frac{A}{n} \mathbb{E}_\sigma \left[\sqrt{\sum_{i,j=1}^n \sigma_i \sigma_j X_i^\top X_j} \right] \leq \frac{A}{n} \sqrt{\sum_{i,j=1}^n \mathbb{E}_\sigma [\sigma_i \sigma_j] X_i^\top X_j} \\ &\leq \frac{A}{n} \sqrt{\sum_{i=1}^n \|X_i\|^2} \leq \frac{AB}{\sqrt{n}} \end{aligned}$$

□

Theorem 2.4. For the perceptron with n input data, i.e. $\mathcal{F}_d = \{x \mapsto \text{sgn}(\theta^\top - \theta_0) : \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}\}$, the growth function is

$$S_{\mathcal{F}_d}(n) = 2 \sum_{k=0}^d \binom{n-1}{k}.$$

Proof. See Anthony, Bartlett, theorem 3.1.

□

2.2 Neural networks

Definition 2.5 (Neural network). A neural network is a function $\Phi : \mathbb{R}^d \rightarrow \mathbb{R}^{N_L}$ given by

$$\Phi(x) = W_L(\rho(W_{L-1}(\rho(\dots \rho(W_1(x))\dots))),$$

where L is the number of layers, N_l the number of parameters (neurons) in layer l for $1 \leq l \leq L$, $\rho : \mathbb{R} \rightarrow \mathbb{R}$ is a (non-linear) function (applied component-wise) and $W_l : \mathbb{R}^{N_{l-1}} \rightarrow \mathbb{R}^{N_l}$ are affine linear maps, i.e. $W_l(x) = A_l x + b_l$, where $A_l \in \mathbb{R}^{N_l \times N_{l-1}}$ and $b_l \in \mathbb{R}^{N_l}$.

- Rademacher complexity for neural network: last example in <http://www.cs.cornell.edu/courses/cs6783/2014fa/lec7.pdf>

2.3 Support vector machines

2.4 Adaptive boosting

The idea of adaptive boosting is to have multiple maybe not so well performing classifiers, and combine them to a single classifier that performs well [? ? ?]. We again consider binary classification, i.e. $\mathcal{Y} = \{-1, 1\}$ and choose K multiple base classifiers $\{f_k\}_{k=1}^K$ with $h_k \in \mathcal{F}$. We then consider the composition classifier $\hat{f}_K(x) = \text{sign} \left\{ \sum_{k=1}^K \alpha_k f_k(x) \right\}$, where $\alpha_k \in \mathbb{R}$ are weights – so the overall classification can be interpreted as a “weighted majority rule”: each datum x gets mapped to a vector $(f_1(x), \dots, f_K(x)) \in \{-1, 1\}^T$ and the weights tell us how to combine these multiple predictions. We now want to study the generalization and therefore estimation error. We

have

$$\begin{aligned}
L_n(\tilde{f}_K) &= \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{\tilde{f}_K(X_i) \neq Y_i\} \\
&= \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{Y_i \sum_{k=1}^K \alpha_k f_k(X_i) \leq 0\} \\
&\leq \sum_{i=1}^n \exp\left(-Y_i \sum_{k=1}^K \alpha_k f_k(X_i)\right) =: \mathcal{L}(\tilde{f}_K),
\end{aligned}$$

where the last line follows from $\mathbb{1}\{z \leq 0\} \leq e^{-z}$. Instead of minimizing $L_n(\tilde{f}_K)$ directly, which can be difficult as it is a non-continuous function, we can minimize the upper bound $\mathcal{L}(\tilde{f}_K)$ and hope that we will find a “similar” minimum. Note that we still consider a minimization over f_1, \dots, f_K as well as $\alpha_1, \dots, \alpha_K$. Since this is computationally heavy we consider an iterative approach, a greedy optimization: Imagine we already tuned f_1, \dots, f_K and $\alpha_1, \dots, \alpha_K$, how do we then choose f_{K+1} and α_{K+1} ? Let us change the writing a bit:

$$\begin{aligned}
\mathcal{L}(\tilde{f}_{K+1}) &= \sum_{i=1}^n \frac{1}{n} \exp\left(-Y_i \underbrace{\sum_{k=1}^K \alpha_k f_k(X_i)}_{=: w_i^{(K)}}\right) \exp(-Y_i \alpha_{K+1} f_{K+1}(X_i)) \\
&= \sum_{i=1}^n w_i^{(K)} \exp(-Y_i \alpha_{K+1} f_{K+1}(X_i)).
\end{aligned}$$

Let $\tilde{w}_i^{(K)} := \frac{w_i^{(K)}}{\sum_{i=1}^n w_i^{(K)}}$, then

$$\begin{aligned}
\frac{\mathcal{L}(\tilde{f}_{K+1})}{\mathcal{L}(\tilde{f}_K)} &= \sum_{i=1}^n \tilde{w}_i^{(K)} \exp(-Y_i \alpha_{K+1} f_{K+1}(X_i)) \\
&= e^{-\alpha_{K+1}} \sum_{i: h_{K+1}(X_i)=Y_i} \tilde{w}_i^{(K)} + e^{\alpha_{K+1}} \sum_{i: h_{K+1}(X_i) \neq Y_i} \tilde{w}_i^{(K)} \\
&= (e^{\alpha_{K+1}} - e^{-\alpha_{K+1}}) \sum_{i=1}^n w_i^{(K)} \mathbb{1}\{f_{K+1}(X_i) \neq Y_i\} + e^{-\alpha_{K+1}}.
\end{aligned}$$

Note that $f_{K+1} \in \mathcal{F}$ only appears in $\epsilon_K := \sum_{i=1}^n w_i^{(K)} \mathbb{1}\{f_{K+1}(X_i) \neq Y_i\}$. Minimization of the above expression gives

$$\alpha_{K+1}^* = \frac{1}{2} \log \frac{1 - \epsilon_K}{\epsilon_K},$$

which is positive if $\epsilon_K < \frac{1}{2}$, which we expect since $\epsilon_K = \frac{1}{2}$ would just correspond to random guessing. Note that

$$w_i^{(K+1)} = w_i^{(K)} e^{-Y_i f_{K+1}(X_i) \alpha_{K+1}^*},$$

which explains the name of “adaptive boosting”. When learning we concentrate on those predictors that have made mistakes in the previous training rounds. We now make a crucial assumption:

$$\epsilon_K < \frac{1}{2} - \gamma, \quad \gamma > 0$$

for every $t = 1, \dots, K$. Note that this assumption depends on choosing a good function class \mathcal{F} . We then have

$$\begin{aligned}
\mathcal{L}(\tilde{f}_{K+1}) &= \mathcal{L}(\tilde{f}_L) \left((e^{\alpha_{K+1}^*} - e^{-\alpha_{K+1}^*}) \epsilon_K + e^{-\alpha_{K+1}^*} \right) \\
&= \mathcal{L}(\tilde{f}_K) 2\sqrt{\epsilon_K(1 - \epsilon_K)} \\
&\leq \mathcal{L}(\tilde{f}_K) 2\sqrt{\left(\frac{1}{2} - \gamma\right) \left(\frac{1}{2} + \gamma\right)} \\
&= \mathcal{L}(\tilde{f}_K) \sqrt{1 - 4\gamma^2} \\
&\leq \mathcal{L}(\tilde{f}_K) e^{-2\gamma^2}
\end{aligned}$$

and therefore

$$\mathcal{L}(\tilde{f}_K) \leq (e^{-2\gamma^2})^K \mathcal{L}(\tilde{f}_0) = e^{-2K\gamma^2}.$$

Since $\mathcal{L}(\tilde{f}_K) \in \{0, \frac{1}{n}, \frac{2}{n}, \dots, 1\}$ we now know that AdaBoost achieves zero empirical loss in finite steps. But does minimizing empirical risk even make sense? Let us say \tilde{f}_K comes from the function class $C(\mathcal{F}, T)$. Note that we have $C(\mathcal{F}, K) \subset C(\mathcal{F}, K+1)$, i.e. the function classes can grow with adding more predictors. That means that the approximation error gets bigger, however at the same time the estimation error could get worse. Vapnik's theorem gave us

$$L(\tilde{f}_K) \leq L_n(\tilde{f}_K) + \mathcal{O}\left(\sqrt{\frac{\text{VC}(C(\mathcal{F}, K)) \log n + \log \frac{1}{\delta}}{n}}\right),$$

where the first summand goes to zero, however, one can show that $\text{VC}(C(\mathcal{F}, K)) \approx K \text{VC}(\mathcal{F})$. So the error should grow sublinearly in T . Still, in practice, people observed that the loss on a test set keeps decreasing even after the empirical risk has gone to zero. To understand why this is the case another approach is necessary that shall be summarized with the following theorem.

Lemma 2.6. *For the Rademacher complexity of convex hulls of function classes we have*

$$\hat{R}_n(\text{conv}(\mathcal{F})) = \hat{R}_n(\mathcal{F}).$$

Remark 2.7. *We have noted before that this is not valid for the VC dimension.*

Theorem 2.8. *Let $\mathcal{Y} = \{-1, 1\}$, $l(y', y) = \mathbb{1}\{y' \neq y\}$. We consider the base class \mathcal{F} with $\text{VC}(\mathcal{F}) < \infty$ and classifiers of the form $\tilde{f}(x) = \text{sign}\{g(x)\}$, where $g \in \text{conv}(\mathcal{F}) := \{g(x) = \sum_{i=1}^N \alpha_i f_i(x), N \in \mathbb{N}, f_1, \dots, f_N \in \mathcal{F}, \alpha_i > 0, \sum_{i=1}^N \alpha_i = 1\}$. Then with probability $\geq 1 - \delta$*

$$L(\tilde{f}) \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{g(X_i)Y_i \leq \eta\} + \frac{2}{\eta} \sqrt{\frac{2 \text{VC}(\mathcal{F}) \log(n+1)}{n}} + \sqrt{\frac{\log \frac{1}{\delta}}{n}}.$$

Proof. Let

$$L(\tilde{f}) = \mathbb{E}[\mathbb{1}\{\tilde{f}(X) \neq Y\}] = \mathbb{E}[\mathbb{1}\{g(X)Y \leq 0\}].$$

Take any $\psi : \mathbb{R} \rightarrow [0, 1]$ that is L_ψ -Lipschitz and for that $\psi(x) \geq \mathbb{1}\{x \leq 0\}$. Then with theorem 1.33

$$\begin{aligned} L(\tilde{f}) &\leq \mathbb{E}[\psi(g(X)Y)] \\ &\leq \frac{1}{n} \sum_{i=1}^n \psi(g(X_i)Y_i) + 2R_n(\psi) + \sqrt{\frac{\log \frac{1}{\delta}}{n}} \end{aligned}$$

with

$$\begin{aligned} R_n(\psi) &= \mathbb{E}_{S_n, \sigma} \left[\sup_{g \in \text{conv}(\mathcal{F})} \frac{1}{n} \sum_{i=1}^n \sigma_i \psi(g(X_i)Y_i) \right] \\ &= 2\mathbb{E}_{S_n, \sigma} \left[\sup_{g \in \text{conv}(\mathcal{F})} \frac{1}{n} \sum_{i=1}^n \sigma_i (\psi(g(X_i)Y_i) - \psi(0)) \right] \\ &\leq 2L_\psi \mathbb{E}_{S_n, \sigma} \left[\sup_{g \in \text{conv}(\mathcal{F})} \frac{1}{n} \sum_{i=1}^n \sigma_i g(X_i)Y_i \right] \\ &= 2L_\psi \mathbb{E}_{S_n, \sigma} \left[\sup_{g \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \sigma_i g(X_i)Y_i \right] \\ &\leq 2L_\psi \sqrt{\frac{2 \text{VC}(\mathcal{F}) \log(n+1)}{n}}, \end{aligned}$$

where we used lemmas 1.49 and ??, theorem 1.46 and the fact that $\sigma_i \stackrel{d}{=} \sigma_i Y_i$. Setting

$$\psi(x) = \begin{cases} 1 & \text{if } x \leq 0 \\ 0 & \text{if } x \geq \eta \\ 1 - \frac{x}{\eta} & \text{in between} \end{cases}$$

yields $L_\psi = \frac{1}{\eta}$ and completes the proof.

- n vs. $n + 1$ – compare lecture 10 from Ilya

□

Remark 2.9. $|g(X_i)Y_i|$ can be interpreted as a margin or a confidence of a prediction $g(X_i)$. Then, the indicator function not only counts the number of wrong classifications, but also the correct classifications that have however not been confident enough (depending on the parameter η).

3 Kernels

Rademacher complexity of Kernel methods

4 Variational inference

- Bayesian approach
- Gaussian processes, Gaussian process regression (see Bartlett, lecture 14)
- PAC Bayes
- check density estimation in Wasserman
- Nonparametric Bayesian methods (see Wasserman)
- Game theoretic context: adversarial model (see also Bartlett, lecture 16)

5 Bibliography

- [1] Bartlett, P. L. and Mendelson, S. (2002). Rademacher and Gaussian complexities: Risk bounds and structural results. *Journal of Machine Learning Research*, 3(Nov):463–482.
- [2] Boucheron, S., Lugosi, G., and Massart, P. (2013). *Concentration inequalities: A nonasymptotic theory of independence*. Oxford university press.
- [3] Bousquet, O., Boucheron, S., and Lugosi, G. (2004). Introduction to statistical learning theory. pages 169–207.
- [4] Cucker, F. and Zhou, D. X. (2007). *Learning theory: an approximation theory viewpoint*, volume 24. Cambridge University Press.
- [5] Devroye, L. and Lugosi, G. (2001). *Combinatorial methods in density estimation*. Springer, New-York.
- [6] Hartmann, C. (2017). *Wahrscheinlichkeitstheorie*. Lecture script, BTU Cottbus-Senftenberg.
- [7] Shalev-Shwartz, S. and Ben-David, S. (2014). *Understanding machine learning: From theory to algorithms*. Cambridge university press.
- [8] Valiant, L. G. (1984). A theory of the learnable. *Communications of the ACM*, 27(11):1134–1142.